

```
In [137]: # importing the libraries
import pandas as pd # for reading dataset
```

```
In [138]: df = pd.read_csv("listings.csv") # Loading the dataset and reading it
```

C:\Users\Sneha\AppData\Local\Temp\ipykernel_27908\2821553885.py:1: DtypeWarning: Columns (68) have mixed types. Specify dtype option on import or set low_memory=False.

```
df = pd.read_csv("listings.csv") # loading the dataset and reading it
```

```
In [139]: df.head() # checking the first 5 rows and cloumns
```

Out[139]:

| | id | listing_url | scrape_id | last_scraped | source | name | description | neighborhood_overview | |
|---|-------|------------------------------------|----------------|--------------|-------------|---|---|---|------------|
| 0 | 92644 | https://www.airbnb.com/rooms/92644 | 20230906022807 | 2023-09-06 | city scrape | Rental unit in Earlsfield · ★4.57 · 1 bedroom ... | The space Hi everyone! I have 2 ro... | NaN | https://aC |
| 1 | 93015 | https://www.airbnb.com/rooms/93015 | 20230906022807 | 2023-09-06 | city scrape | Rental unit in Hammersmith · ★4.82 · 2 bedroom... | Gorgeous 2 bed ground floor apartment with per... | A bit of history about the W14 area: Com... | https://aC |
| 2 | 13913 | https://www.airbnb.com/rooms/13913 | 20230906022807 | 2023-09-06 | city scrape | Rental unit in Islington · ★4.80 · 1 bedroom ... | My bright double bedroom with a large window h... | Finsbury Park is a friendly melting pot commun... | https:// |
| 3 | 45400 | https://www.airbnb.com/rooms/45400 | 20230906022807 | 2023-09-06 | city | Rental unit in London · | Lots of windows and | | https:// |

```
In [129]: # checking columns with unique
df.columns.unique()
```

Out[129]: Index(['id', 'listing_url', 'scrape_id', 'last_scraped', 'source', 'name', 'description', 'neighborhood_overview', 'picture_url', 'host_id', 'host_url', 'host_name', 'host_since', 'host_location', 'host_about', 'host_response_time', 'host_response_rate', 'host_acceptance_rate', 'host_is_superhost', 'host_thumbnail_url', 'host_picture_url', 'host_neighbourhood', 'host_listings_count', 'host_total_listings_count', 'host_verifications', 'host_has_profile_pic', 'host_identity_verified', 'neighbourhood', 'neighbourhood_cleansed', 'neighbourhood_group_cleansed', 'latitude', 'longitude', 'property_type', 'room_type', 'accommodates', 'bathrooms', 'bathrooms_text', 'bedrooms', 'beds', 'amenities', 'price', 'minimum_nights', 'maximum_nights', 'minimum_minimum_nights', 'maximum_minimum_nights', 'minimum_maximum_nights', 'maximum_maximum_nights', 'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', 'calendar_updated', 'has_availability', 'availability_30', 'availability_60', 'availability_90', 'availability_365', 'calendar_last_scraped', 'number_of_reviews', 'number_of_reviews_ltm', 'number_of_reviews_l30d', 'first_review', 'last_review', 'review_scores_rating', 'review_scores_accuracy', 'review_scores_cleanliness', 'review_scores_checkin', 'review_scores_communication', 'review_scores_location', 'review_scores_value', 'license', 'instant_bookable', 'calculated_host_listings_count', 'calculated_host_listings_count_entire_homes', 'calculated_host_listings_count_private_rooms', 'calculated_host_listings_count_shared_rooms', 'reviews_per_month'], dtype='object')

To Keep only necessary columns

```
In [130]: #importing the libraries
import pandas as pd

# selecting columns which are only required
selecting_col = [
    'host_since', 'host_response_time', 'host_response_rate',
    'host_acceptance_rate', 'host_is_superhost', 'host_listings_count',
    'property_type', 'room_type', 'accommodates', 'bathrooms',
    'bedrooms', 'beds', 'amenities', 'price', 'number_of_reviews',
    'latitude', 'longitude', 'neighbourhood_cleansed',
    'review_scores_rating', 'reviews_per_month'
]

# keeping required column only
df1 = df[selecting_col]

#checking the size of the dataset
df1.shape
```

Out[130]: (87946, 20)

```
In [131]: df1.head() #Checking first 5 columns
```

Out[131]:


| | host_since | host_response_time | host_response_rate | host_acceptance_rate | host_is_superhost | host_listings_count | property_type | room_type | a |
|---|------------|--------------------|--------------------|----------------------|-------------------|---------------------|-----------------------------|-----------------|---|
| 0 | 2011-04-10 | NaN | NaN | 100% | f | 1.0 | Private room in rental unit | Private room | |
| 1 | 2011-04-11 | within a few hours | 100% | 25% | f | 1.0 | Entire rental unit | Entire home/apt | |
| 2 | 2009-11-16 | within a few hours | 100% | 88% | f | 3.0 | Private room in rental unit | Private room | |
| 3 | 2009-12-05 | within a day | 100% | 41% | f | 1.0 | Entire rental unit | Entire home/apt | |
| 4 | 2011-04-10 | within a few hours | 90% | 75% | t | 1.0 | Private room in condo | Private room | |

Data visualisation


```
In [132]: import seaborn as sns
import matplotlib.pyplot as plt

# Set the style of seaborn
sns.set(style="whitegrid")

# Pairplot to visualize relationships between numerical columns
sns.pairplot(df1[['accommodates', 'bathrooms', 'bedrooms', 'beds', 'price', 'number_of_reviews']])
plt.show()
```



```
In [ ]: # Countplot of room categories
plt.figure(figsize=(10, 6))
sns.countplot(x='room_type', data=df1)
plt.xticks(rotation=45)
plt.show()
```



```
In [155]: #Matplotlib is an extremely sophisticated Python charting package.
#Seaborn is a high-level interface for providing statistical visuals that is built on Matplotlib.
import matplotlib.pyplot as plt
import seaborn as sns

property_types10 = df1['property_type'].value_counts().nlargest(10).index

# Filtering the dataset for the top 10 property types
df_10 = df1[df1['property_type'].isin(property_types10)]

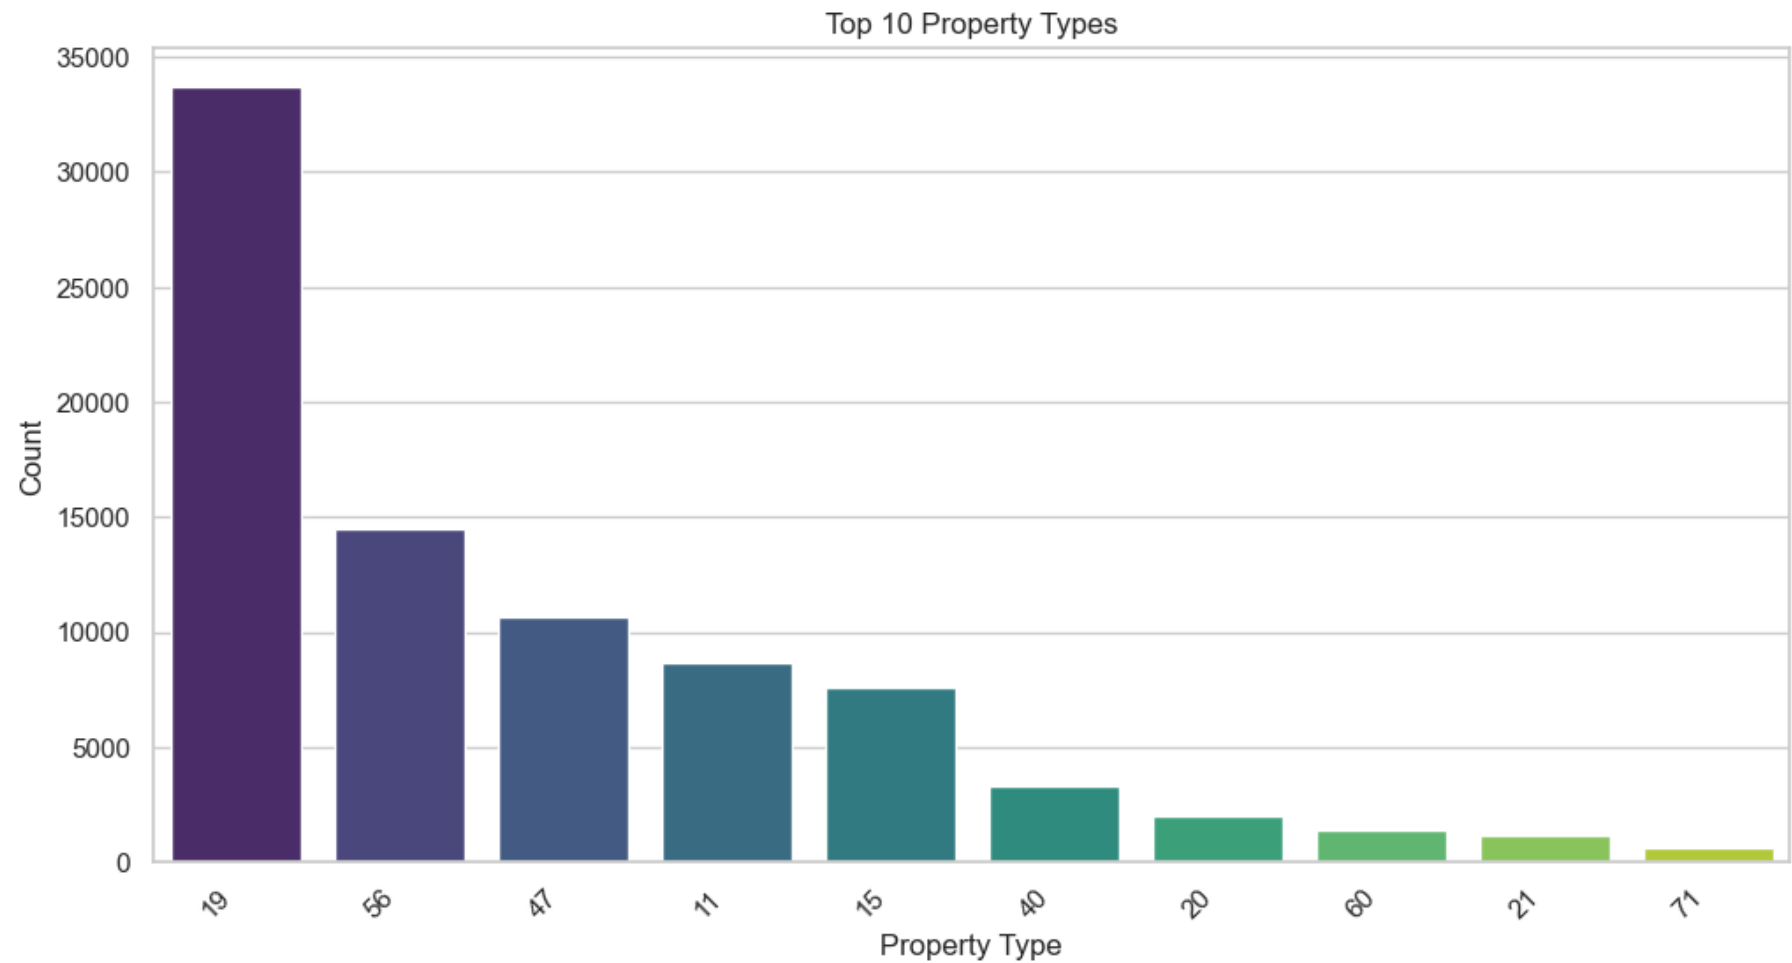
# Chart for top 10 property type
plt.figure(figsize=(12, 6))
sns.countplot(x='property_type', data=df_10, order=property_types10, palette='viridis')
plt.title('Top 10 Property Types')
plt.xlabel('Property Type')
plt.ylabel('Count')
plt.xticks(rotation=45, ha='right')
```

C:\Users\Sneha\AppData\Local\Temp\ipykernel_27908\194507423.py:14: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='property_type', data=df_10, order=property_types10, palette='viridis')
```

```
Out[155]: ([0, 1, 2, 3, 4, 5, 6, 7, 8, 9],
[Text(0, 0, '19'),
Text(1, 0, '56'),
Text(2, 0, '47'),
Text(3, 0, '11'),
Text(4, 0, '15'),
Text(5, 0, '40'),
Text(6, 0, '20'),
Text(7, 0, '60'),
Text(8, 0, '21'),
Text(9, 0, '71')])
```



EDA

Step 1: Checking the datatypes

```
In [140]: # Check data types of each column
d_type = df1.dtypes

# Separating data into numeric and categorical columns
numeric_columns = d_type[d_type != 'object'].index
categorical_columns = d_type[d_type == 'object'].index

# Print the results
print("Numeric Columns:")
print(numeric_columns)

print("\nCategorical Columns:")
print(categorical_columns)
```

Numeric Columns:
Index(['host_listings_count', 'accommodates', 'bathrooms', 'bedrooms', 'beds',
 'number_of_reviews', 'latitude', 'longitude', 'review_scores_rating',
 'reviews_per_month'],
 dtype='object')

Categorical Columns:
Index(['host_since', 'host_response_time', 'host_response_rate',
 'host_acceptance_rate', 'host_is_superhost', 'property_type',
 'room_type', 'amenities', 'price', 'neighbourhood_cleansed'],
 dtype='object')

```
In [141]: from sklearn.preprocessing import LabelEncoder

# Categorical columns to convert
categorical_columns = ['host_since', 'host_response_rate', 'host_acceptance_rate', 'host_is_superhost', 'amenities', 'a

# Initialising LabelEncoder
label_encoder = LabelEncoder()

# Apply label encoding to each categorical column using .loc
for column in categorical_columns:
    df1.loc[:, column] = label_encoder.fit_transform(df1[column])

# Display the updated DataFrame
(df1.head())
```

Out[141]:

| | host_since | host_response_time | host_response_rate | host_acceptance_rate | host_is_superhost | host_listings_count | property_type | room_type | a |
|---|------------|--------------------|--------------------|----------------------|-------------------|---------------------|---------------|-----------|---|
| 0 | 286 | NaN | 95 | 3 | 0 | 1.0 | 56 | 2 | |
| 1 | 287 | within a few hours | 2 | 18 | 0 | 1.0 | 19 | 0 | |
| 2 | 35 | within a few hours | 2 | 87 | 0 | 3.0 | 56 | 2 | |
| 3 | 40 | within a day | 2 | 36 | 0 | 1.0 | 19 | 0 | |
| 4 | 286 | within a few hours | 85 | 73 | 1 | 1.0 | 40 | 2 | |

Checking for Missing values using isnull().sum() it checks for all columns missing values and taking percentage of missing values

```
In [142]: # Check for missing values in the DataFrame
missing_values = df1.isnull().sum()

# Displaying the columns with missing values, if any
Col_missing_val = missing_values[missing_values > 0]
print("Columns with missing values:")
print(Col_missing_val)

# Percentage of missing values in each column
missing_percentage = (missing_values / len(df1)) * 100
print("\nPercentage of missing values in each column:")
print(missing_percentage)
```

Columns with missing values:

| | |
|----------------------|-------|
| host_response_time | 28918 |
| host_listings_count | 5 |
| bathrooms | 87946 |
| bedrooms | 32774 |
| beds | 1134 |
| review_scores_rating | 22158 |
| reviews_per_month | 22158 |

dtype: int64

Percentage of missing values in each column:

| | |
|------------------------|------------|
| host_since | 0.000000 |
| host_response_time | 32.881541 |
| host_response_rate | 0.000000 |
| host_acceptance_rate | 0.000000 |
| host_is_superhost | 0.000000 |
| host_listings_count | 0.005685 |
| property_type | 0.000000 |
| room_type | 0.000000 |
| accommodates | 0.000000 |
| bathrooms | 100.000000 |
| bedrooms | 37.266050 |
| beds | 1.289428 |
| amenities | 0.000000 |
| price | 0.000000 |
| number_of_reviews | 0.000000 |
| latitude | 0.000000 |
| longitude | 0.000000 |
| neighbourhood_cleansed | 0.000000 |
| review_scores_rating | 25.195006 |
| reviews_per_month | 25.195006 |

dtype: float64

```
In [143]: # Dropping the bathroom column since it's showing 100 percent missing values using drop function
df1= df1.drop(['bathrooms'], axis=1)
```

```
In [144]: # The function guarantees that missing values in df1 are replaced with mode values generated from the original Dataset
#df columns. This is a popular approach for dealing with missing data by substituting the most frequent values from ea

df1.loc[:, :] = df1.fillna(df.mode().iloc[0])
df1.isnull().any().sum() # now again checking if there is null values or not
```

Out[144]: 0

Feature Engineering

```
In [145]: # Check summary statistics for columns number_of_reviews,review_scores_rating
print(df1[['number_of_reviews', 'review_scores_rating']].describe())
```

| | number_of_reviews | review_scores_rating |
|-------|-------------------|----------------------|
| count | 87946.000000 | 87946.000000 |
| mean | 17.977236 | 4.698192 |
| std | 42.834975 | 0.672555 |
| min | 0.000000 | 0.000000 |
| 25% | 0.000000 | 4.670000 |
| 50% | 4.000000 | 4.960000 |
| 75% | 16.000000 | 5.000000 |
| max | 1536.000000 | 5.000000 |

To provide descriptive statistics for a DataFrame or Series, utilize Pandas' describe function. The distribution, central tendency, and dispersion of the values in the corresponding columns are all shown by these statistics.

```
In [146]: # Calculating the mean of 'number_of_reviews'
review_threshold = df1['number_of_reviews'].mean()

# Calculating the mean of 'review_scores_rating'
rating_threshold = df1['review_scores_rating'].mean()

# Creating a new column 'Success' based on no of reviews and review score rating
# Updating the 'Success' to 'Bad' using .loc
df1.loc[:, 'Success'] = 'Bad'

# Updating 'Success' to 'Good' based on conditions
condition_reviews = df1['number_of_reviews'] >= review_threshold
condition_rating = df1['review_scores_rating'] >= rating_threshold

# specifying the condition of Good
df1.loc[condition_reviews & condition_rating, 'Success'] = 'Good'

# the updated Dataset
print(df1[['number_of_reviews', 'review_scores_rating', 'Success']])
```

| | number_of_reviews | review_scores_rating | Success |
|-------|-------------------|----------------------|---------|
| 0 | 216 | 4.57 | Bad |
| 1 | 38 | 4.82 | Good |
| 2 | 41 | 4.80 | Good |
| 3 | 94 | 4.80 | Good |
| 4 | 180 | 4.62 | Bad |
| ... | ... | ... | ... |
| 87941 | 0 | 5.00 | Bad |
| 87942 | 0 | 5.00 | Bad |
| 87943 | 0 | 5.00 | Bad |
| 87944 | 0 | 5.00 | Bad |
| 87945 | 0 | 5.00 | Bad |

[87946 rows x 3 columns]

- Determine the average of the columns "number_of_reviews" and "review_scores_rating" in the Dataset df1.
- Assign the value "Bad" to the new column "Success" for every row in df1.
- Set up criteria to determine if the number of reviews ('number_of_reviews') and the review scores ('review_scores_rating') are more than or equal to the computed mean ('reviews_threshold').
- Rows that meet both requirements should have their "Success" column updated to "Good."
- Show the chosen columns of df1 ('number_of_reviews','review_scores_rating', 'Success').

Listings are categorized as "Good" or "Bad" by this code according to factors including the quantity and quality of reviews. The listing is marked as "Good" if both the number of reviews and the review score are more than or equal to the mean; if not, it is marked as "Bad."

```
In [147]: #Checking count of success using value_counts()
success_counts = df1['Success'].value_counts()

# Display the counts
print(success_counts)
```

```
Success
Bad      74298
Good     13648
Name: count, dtype: int64
```

```
In [148]: df1.columns # checking dataset is updated with Success
```

```
Out[148]: Index(['host_since', 'host_response_time', 'host_response_rate',
                'host_acceptance_rate', 'host_is_superhost', 'host_listings_count',
                'property_type', 'room_type', 'accommodates', 'bedrooms', 'beds',
                'amenities', 'price', 'number_of_reviews', 'latitude', 'longitude',
                'neighbourhood_cleansed', 'review_scores_rating', 'reviews_per_month',
                'Success'],
                dtype='object')
```



```

In [149]: #importing all library required
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler

#Defining the Independent and dependent variable and target variable is Success
X = df1.drop('Success', axis=1)
y = df1['Success']

# Initialising the LabelEncoder
label_encoder = LabelEncoder()

# Applying Label Encoding to categorical columns
for col in ['host_response_time', 'property_type', 'room_type', 'neighbourhood_cleansed']:
    df1[col] = label_encoder.fit_transform(df1[col])

# Create features (X) and target variable (y)
X = df1.drop(['Success'], axis=1)
y = df1['Success']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

#Scaling the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)

# Initialize Logistic Regression model
logreg_model = LogisticRegression(random_state=42)

# Fit the model on the training data
logreg_model.fit(X_train_scaled, y_train)

# Predictions on the test set
y_pred = logreg_model.predict(X_test_scaled)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

print("Accuracy:", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", classification_rep)

```

Accuracy: 0.8987492893689596

Confusion Matrix:

```
[[14586  259]
 [ 1522 1223]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Bad | 0.91 | 0.98 | 0.94 | 14845 |
| Good | 0.83 | 0.45 | 0.58 | 2745 |
| accuracy | | | 0.90 | 17590 |
| macro avg | 0.87 | 0.71 | 0.76 | 17590 |
| weighted avg | 0.89 | 0.90 | 0.89 | 17590 |

- Accuracy of the model is 89.87% .This shows how accurate the model is overall in predicting both classifications (Good and Bad).
 - The model's predictions are broken out in great depth in the confusion matrix.The estimated classes are represented by columns, and the actual classifications ('Bad' and 'Good') are represented by rows.
- This model showed that 14586 cases were properly predicted as "Bad," 1223 instances as "Good," 259 instances as "Good" (false positive), and 1522 instances as "Bad" (false negative) were wrongly anticipated.

Precision :The fraction of accurately predicted 'Good' cases among all instances anticipated as 'Good'.

F1-score: The F1-score is the harmonic mean of accuracy and recall. It strikes a balance between accuracy and recall. Support: The number of real occurrences of each class in the dataset. Recall: The recall for 'Good' is 0.45, which is the proportion of properly predicted 'Good' cases out of all actual 'Good' instances.

```
In [151]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Decision Tree Model
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train, y_train)

# Predictions on the test set
y_pred = dt_model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

print(f'Accuracy: {accuracy}')
print(f'Confusion Matrix:\n{conf_matrix}')
print(f'Classification Report:\n{classification_rep}')
```

```
Accuracy: 1.0
Confusion Matrix:
[[14845    0]
 [    0  2745]]
Classification Report:
              precision    recall  f1-score   support

     Bad         1.00        1.00        1.00     14845
     Good         1.00        1.00        1.00       2745

 accuracy                   1.00         17590
 macro avg              1.00         17590
weighted avg              1.00         17590
```

 The Accuracy of the model is 1 which means 100 .

Since the Ratio of Bad and good list is not 50 : 50 ratio.While it may seem ideal to achieve perfect accuracy and performance metrics, doing so increases the risk of overfitting.

 Reason for overfitting

Missing data and noise in data ,lack of regularisation

 Solution

Methods of Regularization:

To penalize large weights in the model, introduce regularization techniques like L1 or L2 regularization. By doing this, the model is kept from fitting the training set too closely.

Cross- validation

Utilizing this methods such as k-fold cross-validation to evaluate the model's performance across several data subsets. This may offer a more thorough assessment of the model's capacity for generalization.

Engineering Features:

Evaluate and improve the features you have. Eliminate any superfluous or pointless features that could be making the model overfit the data's noise.

The Importance of Features used while creating the target variable

The significance of this process lies in the creation of a binary classification ('Good' or 'Bad') based on a combination of two important metrics: the number of reviews and the rating of the review scores. This classification can be useful for identifying successful or popular items, services, or entities within the dataset, among other things. It allows you to divide instances into two groups based on whether they have more or fewer reviews and ratings than the average. This type of categorization can help to simplify analysis and decision-making processes in a variety of applications, including customer satisfaction, product evaluation, and service quality evaluation.

The review rating and review score are significant because they provide concise and quantifiable measures of user satisfaction and product or service quality, which influence consumer decisions and business success. High ratings indicate positive experiences, which attract more customers, whereas low ratings may indicate issues that need to be addressed, affecting reputation and profitability. These metrics provide valuable feedback, guiding businesses to improve customer satisfaction and make informed decisions in order to remain competitive in the market.

Dicussion

we were given the airnb dataset with 75 columns so First i started removing the unnecessary columns and selecting which columns are most important . Then comes data visualisation and seeing which room type has highest value and then comes checking data types and missing values and handling it which plays crucial and important Pre Processing step

Next important step was feature engineering which means create a new column with existing one so i took no of review and review score rating compare there threshold and specify some condition telling it is good or bad

After getting the target variable and i did define dependent and independent variable for model ,scale the features ,applied the logistic and Decision model and predictive the values

so for logistic Regression i got 89 % accuracy and Decision tree 100% and in decision tree it tells that model is overfit